

BINARY IMAGE CLASSIFIER

1. Introduction


This report details the implementation of a binary image classifier (e.g., classifying images as either Cats or Dogs). Building a deep neural network from scratch requires immense computational power and massive datasets. To bypass this limitation, the provided codebase utilizes **Transfer Learning**, specifically adapting a highly efficient pretrained model called **MobileNetV2**. This document elucidates the underlying theory, details the libraries utilized, and provides a systematic, step-by-step breakdown of the implemented code.

2. Theoretical Foundation

Before examining the implementation, it is imperative to establish the three core theoretical concepts governing this architecture.

2.1 Convolutional Neural Networks (CNNs)

Deep learning models tailored for computer vision are known as CNNs. As an image propagates through a CNN, the model applies mathematical filters (convolutions) to detect spatial patterns. Early layers identify rudimentary features such as edges and color gradients, while deeper layers aggregate these to recognize complex, object-specific anatomical features.



2.2 Transfer Learning

Rather than training a CNN from random initialization (which necessitates substantial computational time and data), this approach leverages a model previously trained on millions of images. By "freezing" the pre-trained model's internal layers, its established feature-extraction

capabilities are preserved. Only the final classification layer is replaced and optimized to address the specific binary classification problem.

2.3 Binary Classification & The Sigmoid Function

The objective is to output a singular probability: **0** for the first class (e.g., Cat) and **1** for the second class (e.g., Dog). To achieve this, the final layer of the network utilizes a Sigmoid activation function.

Mathematical Formula:

$$S(x) = 1/(1+e^{-x})$$

This function maps any raw continuous numerical output from the network into a strict probability distribution spanning from **0.0 to 1.0**. A standard decision boundary is applied: outputs strictly below 0.5 are assigned to class 0, whereas outputs of 0.5 or greater are assigned to class 1.

3. Libraries Used

The implementation relies on two foundational libraries within the Python data science ecosystem:

- **TensorFlow (tf)**: An open-source machine learning framework developed by Google. The high-level Keras API (`tf.keras`) is utilized to instantiate, compile, and train the neural networks.
 - **Scikit-Learn (sklearn)**: Utilized strictly for its `metrics` module. It computes the final evaluation metrics (Accuracy, Precision, Recall) via the `classification_report` and `confusion_matrix` functions.
-

4. Code Breakdown and Explanation

Step 1: Loading the Pretrained Model

Python

```
base_model = tf.keras.applications.MobileNetV2(input_shape=(128,128,3), include_top=False,
weights='imagenet')
base_model.trainable = False
```

- **MobileNetV2**: The chosen base architecture, optimized for efficiency.
- **input_shape=(128,128,3)**: Images resized to 128x128 pixels with 3 color channels (RGB).
- **include_top=False**: Removes the original 1,000-class head to allow for custom classification.
- **weights='imagenet'**: Uses weights learned from the ImageNet dataset.
- **base_model.trainable = False**: Freezes the layers to keep pre-learned knowledge intact.

Step 2: Building the Custom Classifier

Python

```
model = tf.keras.Sequential([
    base_model,
    tf.keras.layers.GlobalAveragePooling2D(),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

- **Sequential**: Creates a linear stack of layers.
- **GlobalAveragePooling2D()**: Flattens the 3D data into a 1D array by averaging features.
- **Dense(1, activation='sigmoid')**: A single output neuron for binary probability.

Step 3: Compiling the Model

Python

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

-
- **optimizer='adam'**: The algorithm used to update weights efficiently.
 - **loss='binary_crossentropy'**: The standard loss function for binary tasks.

5. Understanding the Output Metrics

| Metric | Definition in this Context |
|--------|----------------------------|
|--------|----------------------------|

Accuracy Overall proportion of correctly classified images.

Precision Reliability of "Dog" predictions (minimizes false alarms).

Recall Ability to find all actual "Dogs" (minimizes missed detections).

F1-Score Balanced score combining Precision and Recall.

Export to Sheets

6. Conclusion

By leveraging Transfer Learning via MobileNetV2, an efficient binary classification model was successfully constructed. Freezing the pre-trained ImageNet base allowed the network to act as a highly capable feature extractor immediately. This resulted in rapid convergence and high accuracy within just 5 training epochs.

1. Data Acquisition

The code starts by grabbing the raw materials.

Python

```
url = 'https://storage.googleapis.com/mledu-datasets/cats_and_dogs_filtered.zip'  
path = tf.keras.utils.get_file('cats_dogs.zip', origin=url, extract=True)  
data_dir = path.replace('cats_dogs.zip', 'cats_and_dogs_filtered')
```

- **get_file**: This utility downloads the dataset and caches it. The `extract=True` argument tells it to unzip the file immediately.
 - **data_dir**: This simply identifies where the unzipped images are stored on your local machine or cloud environment.
-

2. Loading the Datasets

This section uses `image_dataset_from_directory`, which is the modern standard for loading images in TensorFlow.

Python

```
train_ds = tf.keras.utils.image_dataset_from_directory(  
    f'{data_dir}/train',  
    validation_split=0.2,  
    subset="training",  
    seed=123,  
    image_size=(160, 160)  
)
```

- **validation_split=0.2**: It carves out **20%** of your training data to act as a "practice exam" during training.

- **seed=123**: Ensures that every time you run this, the images are shuffled the same way. Consistency is key for debugging.
 - **image_size=(160, 160)**: This is crucial. Neural networks are picky eaters; every image must be resized to the exact same dimensions before entering the model.
-

3. Building the Model (The "Brain")

This is where the Transfer Learning magic happens.

The Pretrained Base

Python

```
base_model = tf.keras.applications.MobileNetV2(  
    input_shape=(160, 160, 3),  
    include_top=False,  
    weights='imagenet'  
)  
base_model.trainable = False
```

- **MobileNetV2**: A Google-designed model trained on **ImageNet** (1.4 million images). It's famous for being extremely fast and efficient.
- **include_top=False**: ImageNet models are built to classify 1,000 different things. Since you only care about 2 (Cats vs. Dogs), you "decapitate" the model—removing the 1,000-class classifier but keeping the "feature extraction" body.
- **trainable = False**: This "freezes" the weights. You're telling the computer: "Don't try to improve what you've already learned; just use your current knowledge."

The Custom Head

Python

```
model = tf.keras.Sequential([  
    base_model,  
    tf.keras.layers.GlobalAveragePooling2D(),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
)
```

- **GlobalAveragePooling2D**: The base model outputs a complex 3D grid of data. This layer "squashes" that grid into a simple flat list of numbers (a vector).

- `Dense(1, activation='sigmoid')`: The final decision maker. It has one neuron that outputs a value between **0** and **1**. Because of the **Sigmoid** activation, 0 usually represents Cat and 1 represents Dog.
-

4. Compilation and Training

```
Python
model.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy'])
)
model.fit(train_ds, validation_data=val_ds, epochs=1)
```

- **adam**: The industry-standard optimizer that adjusts the model weights to reduce errors.
 - **binary_crossentropy**: The mathematical way we measure how "wrong" the model is for a two-choice problem.
 - **epochs=1**: The model will look at the entire dataset exactly once. In a real-world scenario, you'd likely run this for 5–10 epochs.
-

5. Evaluation and Metrics

Finally, we see how the model performs on the **Test Set** (images it has never seen before).

```
Python
y_true = tf.concat([y for x, y in test_ds], axis=0)
y_pred = (model.predict(test_ds) > 0.5).astype("int32")

print(classification_report(y_true, y_pred))
```

- **tf.concat**: This extracts the true labels (0s and 1s) from the dataset object so we can compare them to our predictions.
- **> 0.5**: We convert the raw probabilities into binary answers. If the model says "0.85 certainty," it becomes a **1** (Dog).
- **classification_report**: This prints a table showing **Precision** (how many "Dog" guesses were actually dogs) and **Recall**(how many of the actual dogs the model managed to find).